



# GPU-Aware KEDA-Based Posture Analyzer - Documentation

This project enables intelligent posture analysis of images sent over MQTT, using GPU-optimized Kubernetes Jobs dynamically scheduled via [KEDA](#). The system is designed to:

- Automatically select the worker node with the lowest GPU load.
- Trigger containerized posture-analysis jobs on that node.
- Log results to a PostgreSQL database.



## Requirements



### Software Dependencies

Install the following using `pip install -r requirements.txt`

```
opencv-python
mediapipe
paho-mqtt
numpy
matplotlib
psycpg2-binary
psutil
prometheus-api-client
keyboard
```



### Environment

- Kubernetes Cluster with KEDA installed
- Prometheus running and scraping all GPU nodes
- Docker with `buildx` and access to Docker Hub (for pushing image)
- Supabase PostgreSQL credentials as environment variables:
  - `SUPABASE_HOST`
  - `SUPABASE_DB`
  - `SUPABASE_USER`
  - `SUPABASE_PASSWORD`
- Optional: `SUPABASE_PORT`, `SUPABASE_SSL`

## Folder and File Overview

File/Folder	Description
<code>build_and_push.py</code>	Main launcher: builds Docker image, applies Kubernetes setup, starts watchers
<code>mqtt_posture_analyzer_with_db.py</code>	Core container logic: receives images, analyzes posture, logs results
<code>gpu_affinity_watcher.py</code>	Selects the worker node with least GPU load and patches KEDA ScaledJob
<code>patch_scaledjob.py</code>	Adds node affinity to <code>posture-job.yaml</code> to force it to run on selected worker
<code>cpu_monitor_and_offload.py</code>	Implements gRPC scaler for KEDA using Prometheus GPU metrics
<code>stop.py</code>	Graceful shutdown script to terminate all posture jobs and clean resources
<code>posture-job.yaml</code>	Base KEDA ScaledJob template
<code>Dockerfile</code>	Builds the posture analysis image
<code>docker-compose.yml</code>	Optional local testing setup
<code>externalscaler.proto</code>	gRPC proto definition for KEDA scaler
<code>externalscaler_pb2.py</code> / <code>externalscaler_pb2_grpc.py</code>	gRPC-generated interfaces for KEDA
<code>requirements.txt</code>	All Python dependencies

## Setup Instructions

### 1. Configure Prometheus

- Ensure your Prometheus instance is running and scraping GPU metrics for all nodes:
- Node 1: `jetson_gpu_usage_percent`
- Node 2: `jetson_orin_gpu_load_percent`

### 2. Build and Push Docker Image

Make sure you're logged in to Docker Hub and then:

```
python3 build_and_push.py
```

This will:

- Stop/remove any existing containers or jobs
- Build a multi-architecture image `shahroz90/posture-analyzer`
- Patch the ScaledJob YAML with node affinity
- Deploy it to the cluster
- Start GPU/CPU watchers and the ESC listener

### 3. MQTT Broker

Set up an MQTT broker at `192.168.1.79` (or modify `mqtt_posture_analyzer_with_db.py` if needed).

- Devices should publish Base64-encoded JPEGs to topics like:
- `images/pi1`
- `images/pi2`

### 4. Supabase Database Setup

Create a PostgreSQL table:

```
CREATE TABLE posture_log (  
  id SERIAL PRIMARY KEY,  
  pi_id TEXT,  
  filename TEXT,  
  received_time TIMESTAMP,  
  analyzed_time TIMESTAMP,  
  neck_angle INTEGER,  
  body_angle INTEGER,  
  posture_status TEXT,  
  landmarks_detected BOOLEAN,  
  processed_by TEXT  
);
```

---

## Runtime Process

Once `build_and_push.py` is executed:

1. **GPU Watcher** (`gpu_affinity_watcher.py`) selects the best GPU node and patches `posture-job.yaml`.
2. **CPU Monitor + ESC Listener** (`cpu_monitor_and_offload.py`) runs a gRPC server for KEDA.
3. KEDA uses the gRPC scaler to decide if GPU usage is low enough to scale.
4. If yes, KEDA schedules a new job based on `posture-job.yaml`.

5. The container runs `mqtt_posture_analyzer_with_db.py`, listens for MQTT images, and logs analysis results.
- 

## Stopping the System

Just run:

```
python3 stop.py
```

It will:

- Kill GPU/CPU watcher scripts
  - Delete all posture-analyzer jobs and pods
  - Untaint the master node (optional)
  - Remove `patched-job.yaml`
- 

## Testing the System

1. Run `python3 build_and_push.py`
  2. Publish an image to `images/pi1` topic over MQTT
  3. Within seconds, the container will:
  4. Receive the image
  5. Analyze posture
  6. Save annotated image to `./analyzed_images/...`
  7. Insert analysis metadata into the PostgreSQL database
- 
-