

# Smart Posture Analyzer using KEDA and K3s - Documentation

This project performs intelligent posture analysis using AI models on containerized jobs, distributed across GPU-equipped nodes based on load. It uses KEDA (Kubernetes Event-Driven Autoscaler) and K3s to manage resource-aware orchestration.

---

## Overview

This system is designed to:

- Receive posture images via MQTT
  - Run posture analysis using MediaPipe in Docker containers
  - Use Prometheus GPU metrics to find the least loaded node
  - Let KEDA scale and dispatch posture-analysis jobs to that node
  - Log the results into a PostgreSQL database (Supabase-hosted)
- 

## Requirements

### Software Dependencies

Install using:

```
pip install -r requirements.txt
```

Dependencies include:

- opencv-python
- mediapipe
- numpy
- paho-mqtt
- psycpg2-binary
- matplotlib
- psutil
- prometheus-api-client
- keyboard

### Environment Setup

- K3s Kubernetes cluster (master + GPU workers)

- KEDA installed in the cluster
- Prometheus monitoring GPU metrics:
  - `jetson_gpu_usage_percent`
  - `jetson_orin_gpu_load_percent`
- Docker (with buildx enabled)
- Docker Hub account (or local registry)
- Supabase PostgreSQL (or alternative DB)

## Folder & File Overview

File/Folder	Description
<code>build_and_push.py</code>	Builds Docker image, sets up taints, launches monitors
<code>mqtt_posture_analyzer_with_db.py</code>	Core script: receives image, runs MediaPipe, logs result
<code>gpu_affinity_watcher.py</code>	Selects GPU node with least load and patches ScaledJob
<code>patch_scaledjob.py</code>	Patches <code>posture-job.yaml</code> with nodeAffinity for selected node
<code>cpu_monitor_and_offload.py</code>	Custom KEDA gRPC scaler using Prometheus metrics
<code>stop.py</code>	Gracefully stops all posture jobs, deletes pods, untaints master
<code>posture-job.yaml</code>	Base KEDA job template (gets patched)
<code>patched-job.yaml</code>	Auto-generated job file with nodeAffinity
<code>externalscaler.proto</code>	gRPC definition for scaler interface
<code>externalscaler_pb2*.py</code>	Auto-generated protobuf/gRPC bindings
<code>Dockerfile</code>	Builds posture analyzer container
<code>docker-compose.yml</code>	Optional: run container locally for test/debug
<code>requirements.txt</code>	Python dependencies list
<code>last_node.txt</code>	Remembers the last selected node to prevent rescheduling

## Setup Instructions

### 1. Configure Prometheus

Ensure Prometheus is scraping GPU usage metrics for all nodes:

- Node 1: `jetson_gpu_usage_percent`
- Node 2: `jetson_orin_gpu_load_percent`

### 2. Build and Push Docker Image

```
python3 build_and_push.py
```

This will:

- Stop/remove existing posture containers
- Build multi-arch image: `shahroz90/posture-analyzer`
- Patch and apply the KEDA ScaledJob
- Start GPU watcher, CPU monitor, and ESC listener

### 3. MQTT Broker

Ensure broker is running at `192.168.1.79` or modify the IP inside `mqtt_posture_analyzer_with_db.py`. Devices should send base64 JPEGs to:

- `images/pi1`
- `images/pi2`, etc.

### 4. Supabase Database Table

```
CREATE TABLE posture_log (  
  id SERIAL PRIMARY KEY,  
  pi_id TEXT,  
  filename TEXT,  
  received_time TIMESTAMP,  
  analyzed_time TIMESTAMP,  
  neck_angle INTEGER,  
  body_angle INTEGER,  
  posture_status TEXT,  
  landmarks_detected BOOLEAN,  
  processed_by TEXT  
);
```

## Runtime Workflow

Once you run:

```
python3 build_and_push.py
```

This happens:

1. **gpu\_affinity\_watcher.py** selects the best GPU node
2. **patch\_scaledjob.py** updates `posture-job.yaml` with nodeAffinity
3. **cpu\_monitor\_and\_offload.py** runs gRPC scaler for KEDA
4. **KEDA** queries Prometheus metrics
5. If eligible, **KEDA triggers posture job** on the chosen node
6. Container runs `mqtt_posture_analyzer_with_db.py`:
7. Subscribes to MQTT topic
8. Decodes image
9. Runs MediaPipe analysis
10. Annotates & saves image (Optional)
11. Inserts metadata into PostgreSQL

---

## Stopping the System

```
python3 stop.py
```

This will:

- Kill GPU/CPU watcher processes
- Delete all Kubernetes pods and jobs
- Untaint master node if tainted
- Delete `patched-job.yaml`

---

## Testing the System

1. Run `python3 build_and_push.py`
  2. Publish a test image to topic: `images/pi1`
  3. Posture-analyzer job should trigger on selected node
  4. (Optional) Annotated image saved to `./analyzed_images/analyzed_images_from_pi1/`
  5. Result logged in PostgreSQL
-

## Summary

- Uses KEDA to autoscale posture jobs based on GPU availability
- Patches job templates dynamically to enforce nodeAffinity
- Combines Docker, Prometheus, and MQTT in a smart edge AI pipeline
- Ensures posture analysis is always run on best-performing node